

Gitlab-CI: Setting up Continuous Integration for a Gitlab Project

Including Creating a Runner

Duncan C. White,
d.white@imperial.ac.uk

Dept of Computing,
Imperial College London

1st Feb 2016

- Gitlab-CI is a new Gitlab feature, enabling you to setup automatic actions (eg testing, integration with other components, package building etc) on your Gitlab projects, that run every time you push changes up.

- Gitlab-CI is a new Gitlab feature, enabling you to setup automatic actions (eg testing, integration with other components, package building etc) on your Gitlab projects, that run every time you push changes up.
- You can add CI to an existing Gitlab project, but this tutorial creates a new Gitlab project, containing a simplified version of a C-based program that flattens nested mailing lists.

- Gitlab-CI is a new Gitlab feature, enabling you to setup automatic actions (eg testing, integration with other components, package building etc) on your Gitlab projects, that run every time you push changes up.
- You can add CI to an existing Gitlab project, but this tutorial creates a new Gitlab project, containing a simplified version of a C-based program that flattens nested mailing lists.
- Log in to the `gitlab.doc.ic.ac.uk` web interface and create a new project called **mini-list-flattening**.

- Gitlab-CI is a new Gitlab feature, enabling you to setup automatic actions (eg testing, integration with other components, package building etc) on your Gitlab projects, that run every time you push changes up.
- You can add CI to an existing Gitlab project, but this tutorial creates a new Gitlab project, containing a simplified version of a C-based program that flattens nested mailing lists.
- Log in to the `gitlab.doc.ic.ac.uk` web interface and create a new project called **mini-list-flattening**.
- Then populate it as follows:

```
pushd /tmp
wget http://www.doc.ic.ac.uk/~dcw/mini-list-flattening.tgz
tar xzf mini-list-flattening.tgz
cd mini-list-flattening
```

- Then follow the "Existing folder" instructions on the newly created Gitlab project page. For me, these were:

```
git init
git remote add origin git@gitlab.doc.ic.ac.uk:dcw/mini-list-flattening.git
git add .
git commit -m "first commit"
git push -u origin master
```

- In Gitlab, click again on the **Project** button and you should see the results of the first commit. On **Project Settings**, check that the **Builds** feature is enabled, and, further down, that your project has a **CI Token**. You'll need to copy this token later.

- In Gitlab, click again on the **Project** button and you should see the results of the first commit. On **Project Settings**, check that the **Builds** feature is enabled, and, further down, that your project has a **CI Token**. You'll need to copy this token later.
- In your new **mini-list-flattening** repo directory, look around and see what the code does. In particular, read the README and do what it says to compile and run the test program.

- In Gitlab, click again on the **Project** button and you should see the results of the first commit. On **Project Settings**, check that the **Builds** feature is enabled, and, further down, that your project has a **CI Token**. You'll need to copy this token later.
- In your new **mini-list-flattening** repo directory, look around and see what the code does. In particular, read the README and do what it says to compile and run the test program.
- No Gitlab project will use CI unless you set up a YAML file called **.gitlab-ci.yml** defining the actions to run. You will spend a lot of time editing/committing/pushing this file, until it works.

- In Gitlab, click again on the **Project** button and you should see the results of the first commit. On **Project Settings**, check that the **Builds** feature is enabled, and, further down, that your project has a **CI Token**. You'll need to copy this token later.
- In your new **mini-list-flattening** repo directory, look around and see what the code does. In particular, read the README and do what it says to compile and run the test program.
- No Gitlab project will use CI unless you set up a YAML file called **.gitlab-ci.yml** defining the actions to run. You will spend a lot of time editing/committing/pushing this file, until it works.
- After a lot of failures, the first roughly correct version read:

```
before_script:
  - sudo apt-get update -qq && sudo apt-get install -y -qq gcc make
  - which gcc

runtests:
  script:
    # compile it up
    - export TOOLDIR=$HOME/c-tools
    - export ARCH=x86_64
    - make
    # and run the tests..
    - make test
```

- In the YAML file, there can be any number of sections. The **before_script** section is special, and means *do each command in the list at the beginning of every test build*.

- In the YAML file, there can be any number of sections. The **before_script** section is special, and means *do each command in the list at the beginning of every test build*.
- The second section can be named whatever you like, here **runtests** was my choice. The **script** tag means *run a sequence of commands*: first we compile the software, then we test it:

```
runtests:
  script:
    # compile it up
    - export TOOLDIR=$HOME/c-tools
    - export ARCH=x86_64
    - make
    # and run the tests..
    - make test
```

- The commands are exactly those (bash syntax) commands that the README file told us to use to build the program and run the tests. They assume that a \$HOME/c-tools directory exists, containing a few useful modules that I use in most C projects.
- Create yourself a **.gitlab-ci.yml** file with the full contents from page 3, then `git add it`, `git commit it` and `git push it` up.

- When you push this file up to your remote repo on Gitlab, Gitlab will automatically enable CI facilities on the project.
- In the Gitlab UI, there's a **Builds** menu item, click on it and you will see that it attempted to run a build, initially this will be marked as Pending.
- Click on the Pending Build and you'll see that it's Pending because you haven't yet created and registered a **Runner**.
- A **Runner** is a special test machine, belonging to you and running special software, that Gitlab-CI will use to run your actions. We'll see how to set up and customize the Runner later in this tutorial. In particular, we'll need to ensure that the `$HOME/c-tools` directory exists with the right contents.

- When you push this file up to your remote repo on Gitlab, Gitlab will automatically enable CI facilities on the project.
- In the Gitlab UI, there's a **Builds** menu item, click on it and you will see that it attempted to run a build, initially this will be marked as Pending.
- Click on the Pending Build and you'll see that it's Pending because you haven't yet created and registered a **Runner**.
- A **Runner** is a special test machine, belonging to you and running special software, that Gitlab-CI will use to run your actions. We'll see how to set up and customize the Runner later in this tutorial. In particular, we'll need to ensure that the `$HOME/c-tools` directory exists with the right contents.
- Note: If you see the Build marked as Failed, not Pending, at this stage, it probably means there are syntax errors in the YML file - these are displayed on the Build page. Fix the YML file, commit it and push it up and check the Build status again.

- The Runner machine could be a physical machine, or a VM, or a docker container. We're going to create a VM on the DoC private cloud.
- To create the Runner VM, log onto the `runner-creator.doc.ic.ac.uk` web interface and create a new VM based on the Featured **Non-CSG Ubuntu 14.04 30GB disk** template, which comes with college authentication and local root console access, but which does not mount DoC/College home dirs, or run the CSG maintenance system.
- All you need to do is enter your DoC login, password, and a short vm name. Please note, the vm name is only a name in cloudstack and on the vm (`/etc/hosts` and `/etc/hostname`). The name is not registered in the DNS database, and so will not resolve.

- There is quite a bit of asynchronous communication happening in the background. If you get a 502 NGINX Bad Gateway error, please check cloudstack (<https://cloudstack.doc.ic.ac.uk>). There is a good chance that the front end timed out before the application was completed. Your vm was probably created.
- You can ssh into the vm either by it's IP address, or by going to cloud-vm-⟨subnet⟩-⟨last byte⟩ where subnet is the third byte, and last byte is the fourth byte of the address.
Example, 146.169.46.65 would be cloud-vm-46-65.
- Check that you have sudo access:

```
id
```

You should see a reference to the sudo group.

Then you can become su:

```
sudo -s
```

- Do the rest of the setup as root via the ssh session.
- The package gitlab-ci-multi-runner is installed in the template.
To connect Gitlab and our new runner VM together, register the runner:

```
gitlab-ci-multi-runner register
```


- Do the rest of the setup as root via the ssh session.
- The package gitlab-ci-multi-runner is installed in the template. To connect Gitlab and our new runner VM together, register the runner:

```
gitlab-ci-multi-runner register
```

- This asks us a few questions, first we enter the name of our Gitlab server's CI endpoint: `https://gitlab.doc.ic.ac.uk/ci`
- Then we copy and paste in our repo's Gitlab CI token (from the Gitlab **Project Settings** page).
- Then we enter a name for the runner - I chose `my-gitlab-ci-runner-ubuntu14.04-cvm` - and enter zero or more symbolic tags - I entered none.
- Finally we choose the *Shell Executor* - I entered `shell`.

- As soon as we have finished registering the multi-runner, go back to the Gitlab web interface, and check the Build status. Remember, up to now, the Build status has been *Pending*, because with no runner Gitlab can't run any tests.

- As soon as we have finished registering the multi-runner, go back to the Gitlab web interface, and check the Build status. Remember, up to now, the Build status has been *Pending*, because with no runner Gitlab can't run any tests.
- But now, the Build status should change in a few seconds from Pending to Failed, and the log panel (black background) should show the commands it ran and the results.
- You should see the Build successfully clone the repo on the runner VM, then fail at the first `sudo apt-get` command.

- As soon as we have finished registering the multi-runner, go back to the Gitlab web interface, and check the Build status. Remember, up to now, the Build status has been *Pending*, because with no runner Gitlab can't run any tests.
- But now, the Build status should change in a few seconds from Pending to Failed, and the log panel (black background) should show the commands it ran and the results.
- You should see the Build successfully clone the repo on the runner VM, then fail at the first `sudo apt-get` command.
- Why? After some investigation, I realised that the build runs as a local **gitlab-runner** user on the runner VM, and that user cannot use `sudo` by default. To discover this, I added `whoami` as an extra command in the **before_script** section in the YML file, then re-committed and re-pushed it:

```
vi .gitlab-ci.yml [added "- whoami" as 1st command in before_script list]
git commit .gitlab-ci.yml
git push
```

- So, to allow our runner to use sudo, add the **gitlab-runner** user to the sudo group, in the root ssh session:

```
usermod -G sudo gitlab-runner
```

- Now, click **Retry Build**. If you got it right, you should see that the runner VM installs gcc and make, as the **before_script** section told it to.

- So, to allow our runner to use sudo, add the **gitlab-runner** user to the sudo group, in the root ssh session:

```
usermod -G sudo gitlab-runner
```

- Now, click **Retry Build**. If you got it right, you should see that the runner VM installs gcc and make, as the **before_script** section told it to.
- Of course, every change made to the runner VM - either by the **before_script** section, or done manually as root on the VM, persists forever. Hence, once we've successfully installed gcc and make, we probably don't want to leave the apt-get commands live, because they run every time and slow things down. So comment most of the **before_script** section out (by another vi; commit and push sequence):

```
before_script:
  #- whoami
  #- sudo apt-get update -qq && sudo apt-get install -y -qq gcc make
  - which gcc
```

- Of course, we could have avoided the **before_script** section, and run the apt-get commands manually in the root ssh shell.

- After completing the **before_script** section, the build process attempted the **runtests** section, which reads:

```
runtests:
  script:
    # compile it up
    - export TOOLDIR=$HOME/c-tools
    - export ARCH=x86_64
    - make
    # and run the tests..
    - make test
```

- After completing the **before_script** section, the build process attempted the **runtests** section, which reads:

```
runtests:
  script:
    # compile it up
    - export TOOLDIR=$HOME/c-tools
    - export ARCH=x86_64
    - make
    # and run the tests..
    - make test
```

- When I did this, the Build process cd'd into the correct directory, set the above environment variables, and then ran `make`.

- After completing the **before_script** section, the build process attempted the **runtests** section, which reads:

```
runtests:
  script:
    # compile it up
    - export TOOLDIR=$HOME/c-tools
    - export ARCH=x86_64
    - make
    # and run the tests..
    - make test
```

- When I did this, the Build process cd'd into the correct directory, set the above environment variables, and then ran `make`.
- At this point, of course, `$HOME/c-tools` did not exist on the runner VM, so the `make` failed to find **mem.h** anywhere.
- We could add rules to the **before_script** section to fetch a **c-tools** tarball from somewhere and extract it, but it's simpler to do this from the root ssh session.
- First, as you on a DoC workstation, build a **c-tools.tgz** tarball containing your `~/c-tools` directory:

```
cd
tar czf /tmp/c-tools.tgz c-tools
```

- Then copy the tarball to the runner VM. I did:

```
scp /tmp/c-tools.tgz cloud-vm-46-64:/tmp
```

- Then, in the runner VM root session:

```
cd /home/gitlab-runner  
tar xf /tmp/c-tools.tgz
```

- Then copy the tarball to the runner VM. I did:

```
scp /tmp/c-tools.tgz cloud-vm-46-64:/tmp
```

- Then, in the runner VM root session:

```
cd /home/gitlab-runner
tar xf /tmp/c-tools.tgz
```

- Now, click **Retry Build** - you should see that the runner VM successfully compiles and links **testmld**, and then runs it, producing the output:

```
basic members: { a,b,c,d,e,ldk,dcw,gnb, }
lists initially:
```

```
two: basic: { a,d,ldk,dcw, }
one: basic: { b,c, }, non-basic: { two, }
three: basic: { e, }, non-basic: { two,one, }
```

```
T allbasic( one: b,c, nonbasic two, ): is false: ok
T allbasic( two: a,d,ldk,dcw, ): is true: ok
T allbasic( three: e, nonbasic two,one, ): is false: ok
...
T allbasic( one: a,b,c,d,ldk,dcw, ): is true: ok
T allbasic( two: a,d,ldk,dcw, ): is true: ok
T allbasic( three: a,b,c,d,e,ldk,dcw, ): is true: ok
```

Build succeeded.

- Note that the Build process fetches the output from the runner and presents it to us unaltered, recording it for posterity.

- Our Build is now - for the first time - successful. Clicking back on Gitlab's **Builds** menu item shows the whole Build history, many failures plus one successful build.
- Now, every time you make any change to your repo, and push it up to Gitlab, another Build will automatically happen, and a few seconds later Gitlab's **Builds** section will show whether the new version ran the tests successfully - try this a few times.
- You may well want to invest a little time wrapping your test runs in some small test harness that summarises (as Perl's **prove** does) your test runs, in order to make it clearer how many tests there were, and how many failed.
- Some languages will have their own test framework you should use, but in our case you will notice that the output comprises *informational messages* interleaved with *test success/failure messages*, the latter marked with a "**T**" prefix.
- As a first step, change the **Makefile** test invocation to:

```
test:  testmld
      ./testmld | grep '^T '
```

- Then `git commit Makefile` and `git push` the change up. In a few seconds, the latest Build output will only show:

```
T allbasic( one: b,c, nonbasic two, ): is false: ok
T allbasic( two: a,d,ldk,dcw, ): is true: ok
T allbasic( three: e, nonbasic two,one, ): is false: ok
T allbasic( one: a,b,c,d,ldk,dcw, ): is true: ok
T allbasic( two: a,d,ldk,dcw, ): is true: ok
T allbasic( three: e, nonbasic two,one, ): is false: ok
T allbasic( one: a,b,c,d,ldk,dcw, ): is true: ok
T allbasic( two: a,d,ldk,dcw, ): is true: ok
T allbasic( three: a,d,e,ldk,dcw, nonbasic one, ): is false: ok
T allbasic( one: a,b,c,d,ldk,dcw, ): is true: ok
T allbasic( two: a,d,ldk,dcw, ): is true: ok
T allbasic( three: a,b,c,d,e,ldk,dcw, ): is true: ok
```

- A simple Perl script can be used to produce more **prove** like output; I've provided one for you - see the **summarisetests** script already present in **mini-list-flattener**. To use it, change the Makefile invocation of **testmld** to read:

```
test: testmld
    ./testmld | ./summarisetests
```

- Test it yourself via `make test` to familiarise yourself with the summarised output.
- Then `git commit Makefile` and `git push` the change up.

- In a few seconds, the latest Build output will be summarised to:

```
12 tests: all 12 pass
```

```
passes:
```

```
allbasic( one: b,c, nonbasic two, ): is false
allbasic( two: a,d,ldk,dcw, ): is true
allbasic( three: e, nonbasic two,one, ): is false
allbasic( one: a,b,c,d,ldk,dcw, ): is true
allbasic( two: a,d,ldk,dcw, ): is true
allbasic( three: e, nonbasic two,one, ): is false
allbasic( one: a,b,c,d,ldk,dcw, ): is true
allbasic( two: a,d,ldk,dcw, ): is true
allbasic( three: a,d,e,ldk,dcw, nonbasic one, ): is false
allbasic( one: a,b,c,d,ldk,dcw, ): is true
allbasic( two: a,d,ldk,dcw, ): is true
allbasic( three: a,b,c,d,e,ldk,dcw, ): is true
```

```
Build succeeded.
```

- In a few seconds, the latest Build output will be summarised to:

```
12 tests: all 12 pass
passes:
  allbasic( one: b,c, nonbasic two, ): is false
  allbasic( two: a,d,ldk,dcw, ): is true
  allbasic( three: e, nonbasic two,one, ): is false
  allbasic( one: a,b,c,d,ldk,dcw, ): is true
  allbasic( two: a,d,ldk,dcw, ): is true
  allbasic( three: e, nonbasic two,one, ): is false
  allbasic( one: a,b,c,d,ldk,dcw, ): is true
  allbasic( two: a,d,ldk,dcw, ): is true
  allbasic( three: a,d,e,ldk,dcw, nonbasic one, ): is false
  allbasic( one: a,b,c,d,ldk,dcw, ): is true
  allbasic( two: a,d,ldk,dcw, ): is true
  allbasic( three: a,b,c,d,e,ldk,dcw, ): is true
```

Build succeeded.

- That's enough for now. In these notes, you've seen how to set up a fresh Gitlab project repository to use **Gitlab-CI** to do automatic testing.

- In a few seconds, the latest Build output will be summarised to:

```
12 tests: all 12 pass
passes:
  allbasic( one: b,c, nonbasic two, ): is false
  allbasic( two: a,d,ldk,dcw, ): is true
  allbasic( three: e, nonbasic two,one, ): is false
  allbasic( one: a,b,c,d,ldk,dcw, ): is true
  allbasic( two: a,d,ldk,dcw, ): is true
  allbasic( three: e, nonbasic two,one, ): is false
  allbasic( one: a,b,c,d,ldk,dcw, ): is true
  allbasic( two: a,d,ldk,dcw, ): is true
  allbasic( three: a,d,e,ldk,dcw, nonbasic one, ): is false
  allbasic( one: a,b,c,d,ldk,dcw, ): is true
  allbasic( two: a,d,ldk,dcw, ): is true
  allbasic( three: a,b,c,d,e,ldk,dcw, ): is true
```

Build succeeded.

- That's enough for now. In these notes, you've seen how to set up a fresh Gitlab project repository to use **Gitlab-CI** to do automatic testing.
- Note that **gitlab-ci-multi-runner** can be used for testing several of your Gitlab projects. Set up Gitlab-CI for a second Gitlab project repo (as before) and then, on your existing runner VM, just rerun the registration - using the second project's CI Token.